

CS 569

Selected Topics in Software Engineering:
Program Analysis & Evaluation

Random Testing

Oregon State University, Winter 2024

Random Testing (Fuzzing)

- Idea: feed random inputs to a program
- Observe whether it behaves “correctly”
 - Output is correct as per specification
 - Does not crash (universal test oracle)
- Special case of mutation analysis

The Infinite Monkey Theorem

Given enough time, a hypothetical monkey typing at random would, as part of its output, almost surely produce one of Shakespeare's plays (or any other text).



Source: https://en.wikipedia.org/wiki/Infinite_monkey_theorem

The First Fuzzing Study

- Conducted by Barton Miller at University of Wisconsin
- 1990: Command-line fuzzer, testing reliability of UNIX programs
 - Bombards utilities with random data
- 1995: Expanded to GUI-based applications (X Windows), network protocols, and system library APIs
- Later: Command-line and GUI-based Windows and OS X apps

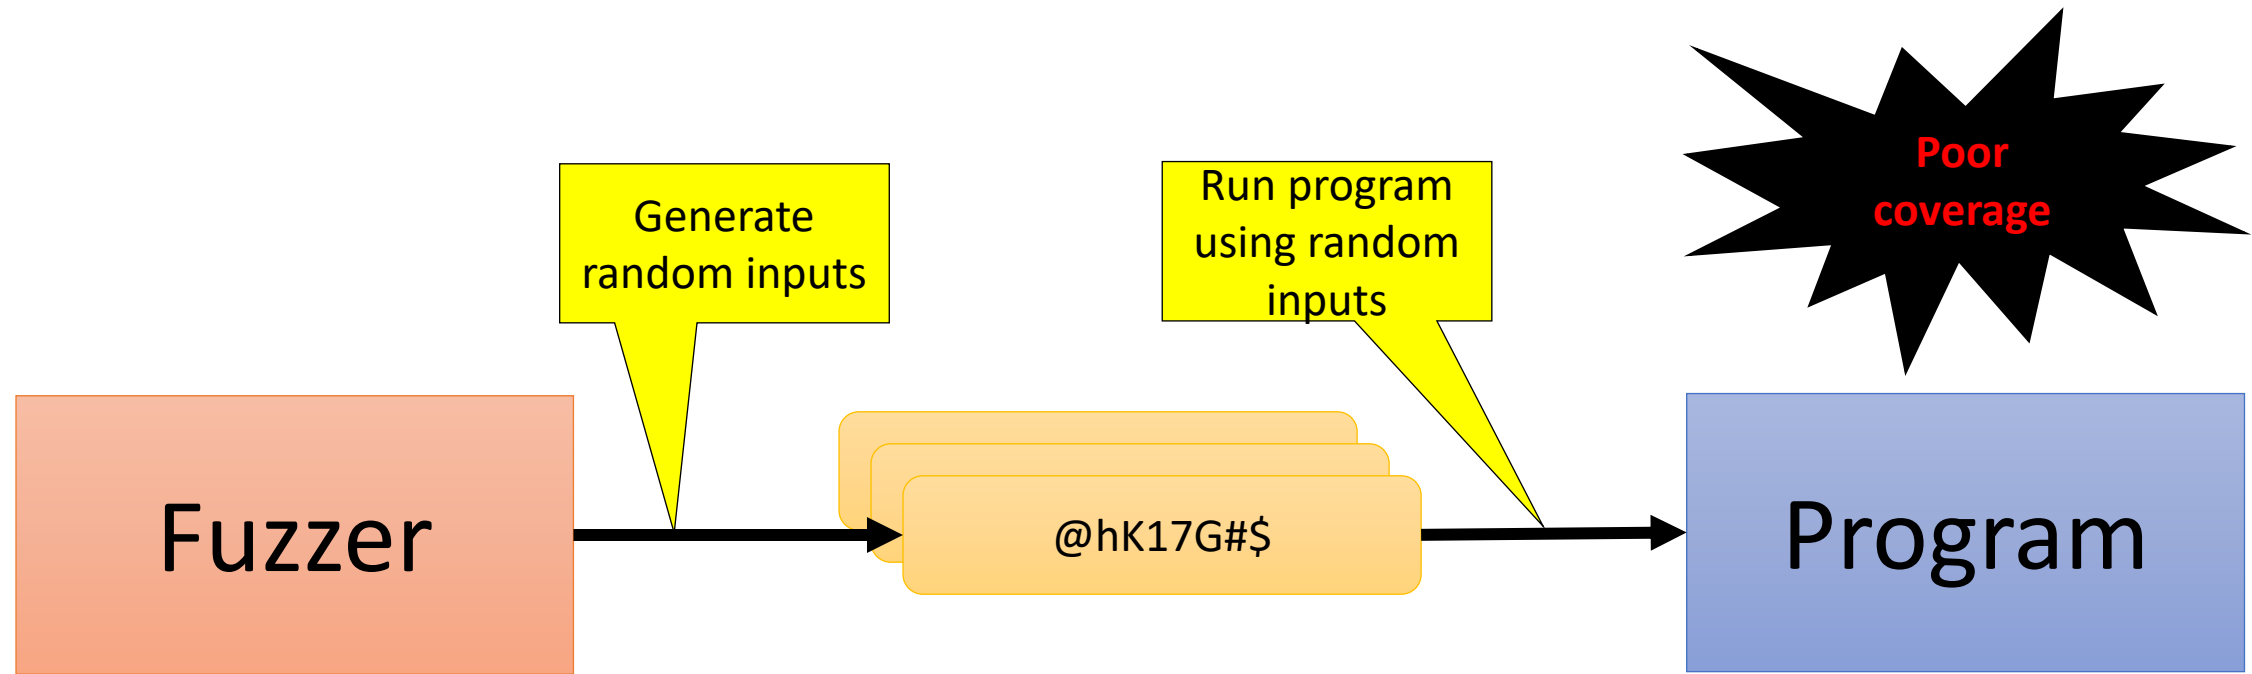
Fuzzing UNIX Utilities: Aftermath

- 1990: Command-line fuzzer, testing reliability of 88 UNIX programs -> caused 25-33% of UNIX utility programs to crash or hang
- 1995: Expanded to GUI-based applications (X Windows), network protocols, and system library APIs -> systems got better but not much
- “Even worse is that many of the bugs reported in 1990 are still present in code releases of 1995”

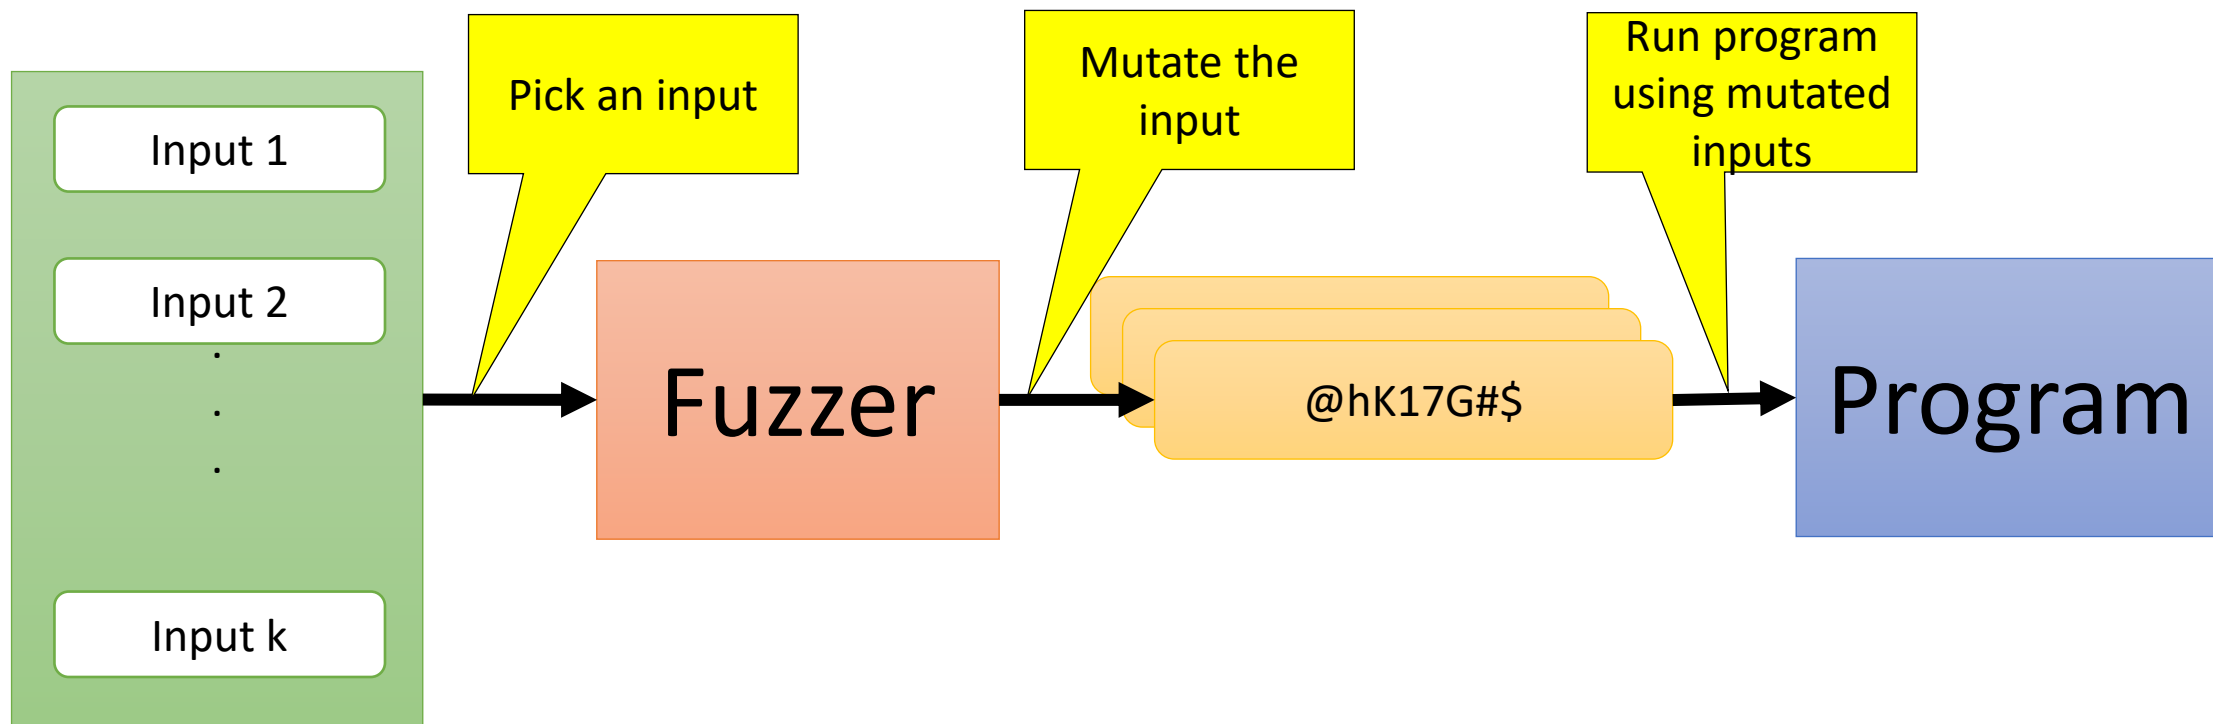
A Silver Lining: Security Bugs

- `gets()` function in C has no parameter limiting input length
=> programmer must make assumptions about the structure of input
- Causes reliability issues and security breaches
 - Second most common cause of errors in 1995 study
- Solution: use `fgets()`, which includes argument to limit the maximum length of input data

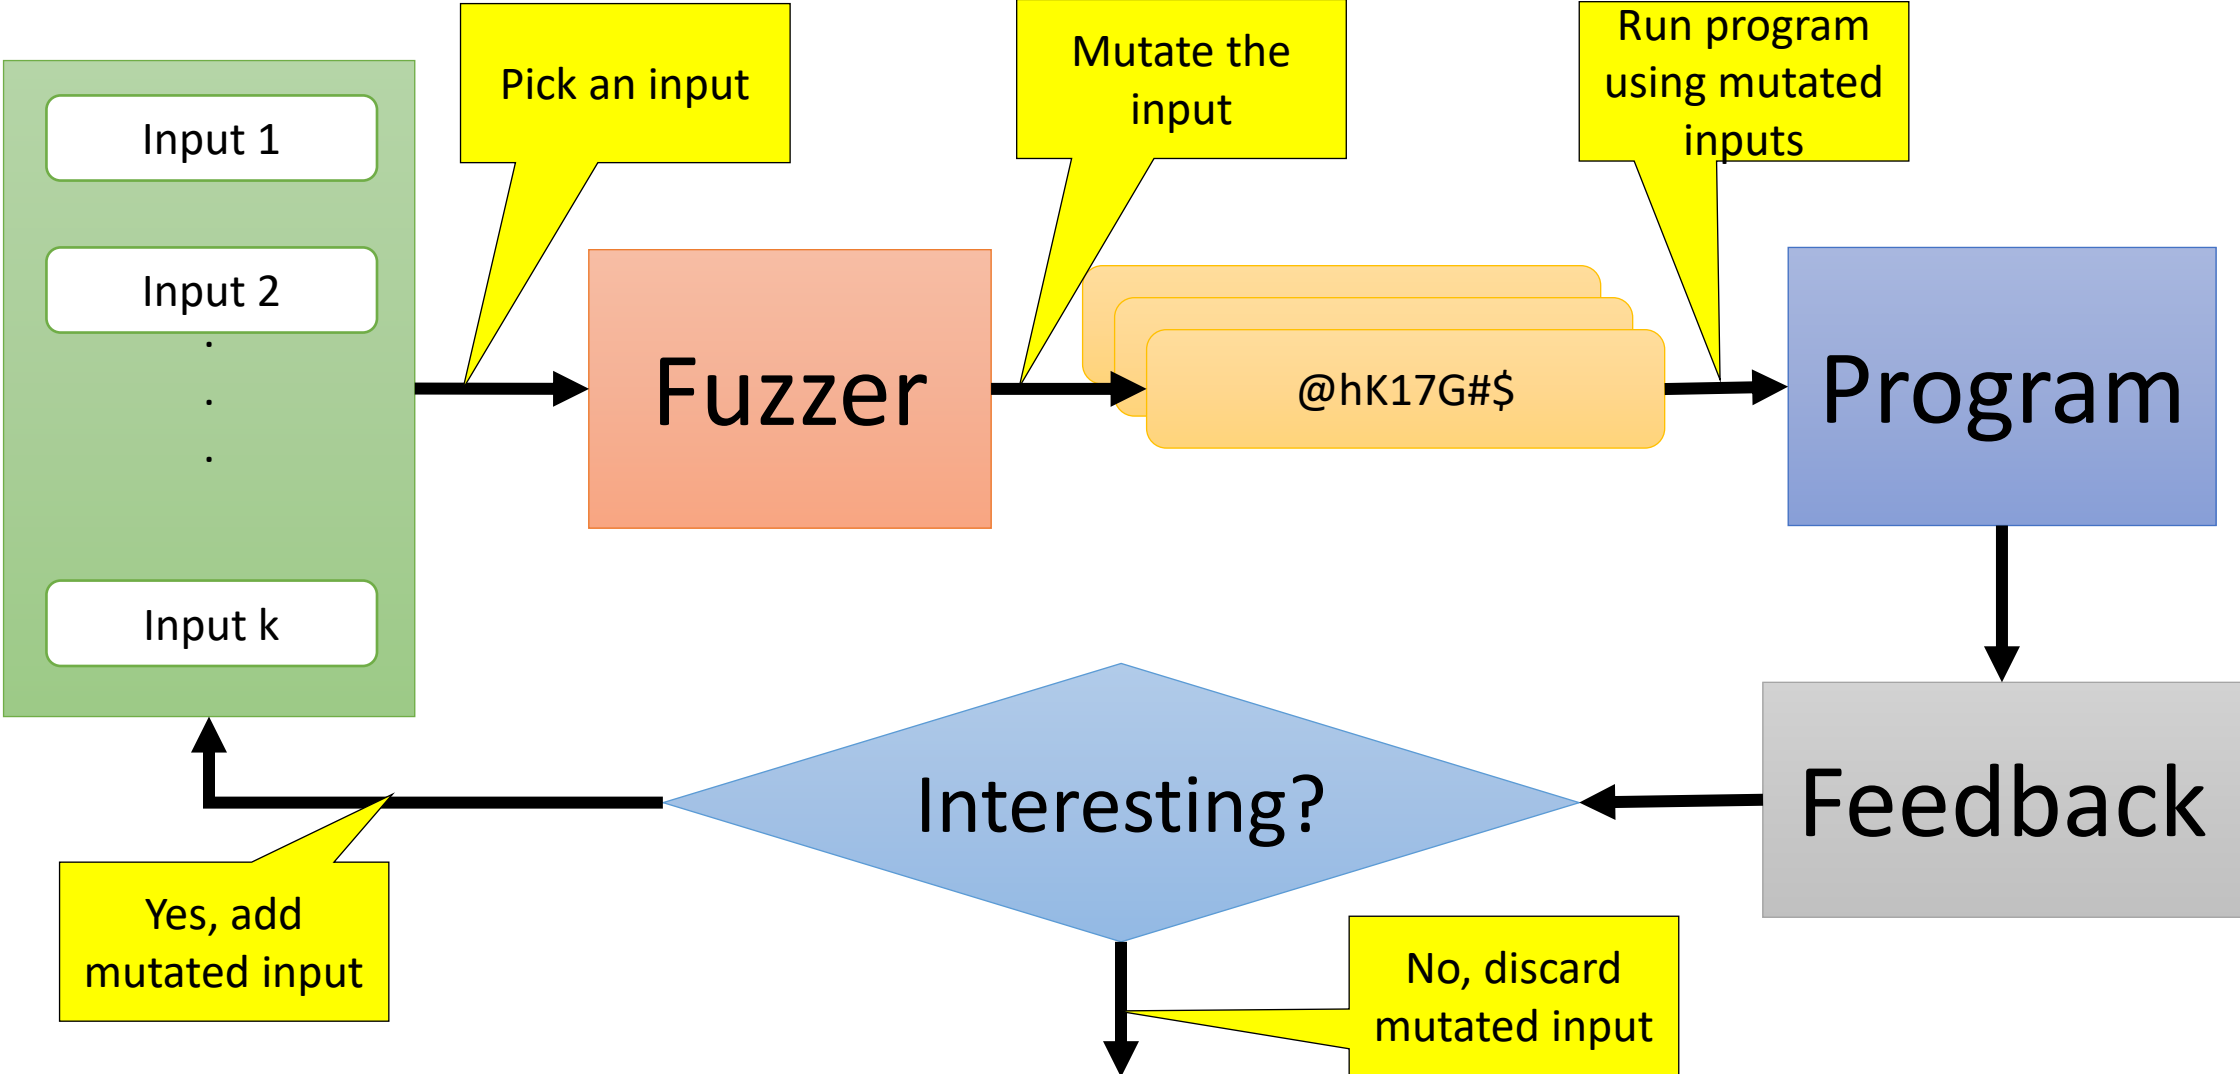
First Generation Fuzzer



Second Generation Fuzzer



Third Generation Fuzzer



What Kinds of Bugs can Fuzzing Find?

- Memory errors
 - Spatial (out-of-bound access) and Temporal (use-after-free)
- Other undefined behaviors
 - Integer overflow, null dereference, divide by zero, uninitialized read, ...
- Assertion violations
- Infinite loops (using timeout)
- Concurrency bugs
 - Data races, deadlocks, ...

Random Testing: Pros and Cons


- Pros

- Easy to implement
- Provably good coverage given enough tests
- Can work with programs of any format
- Appealing to find security vulnerabilities

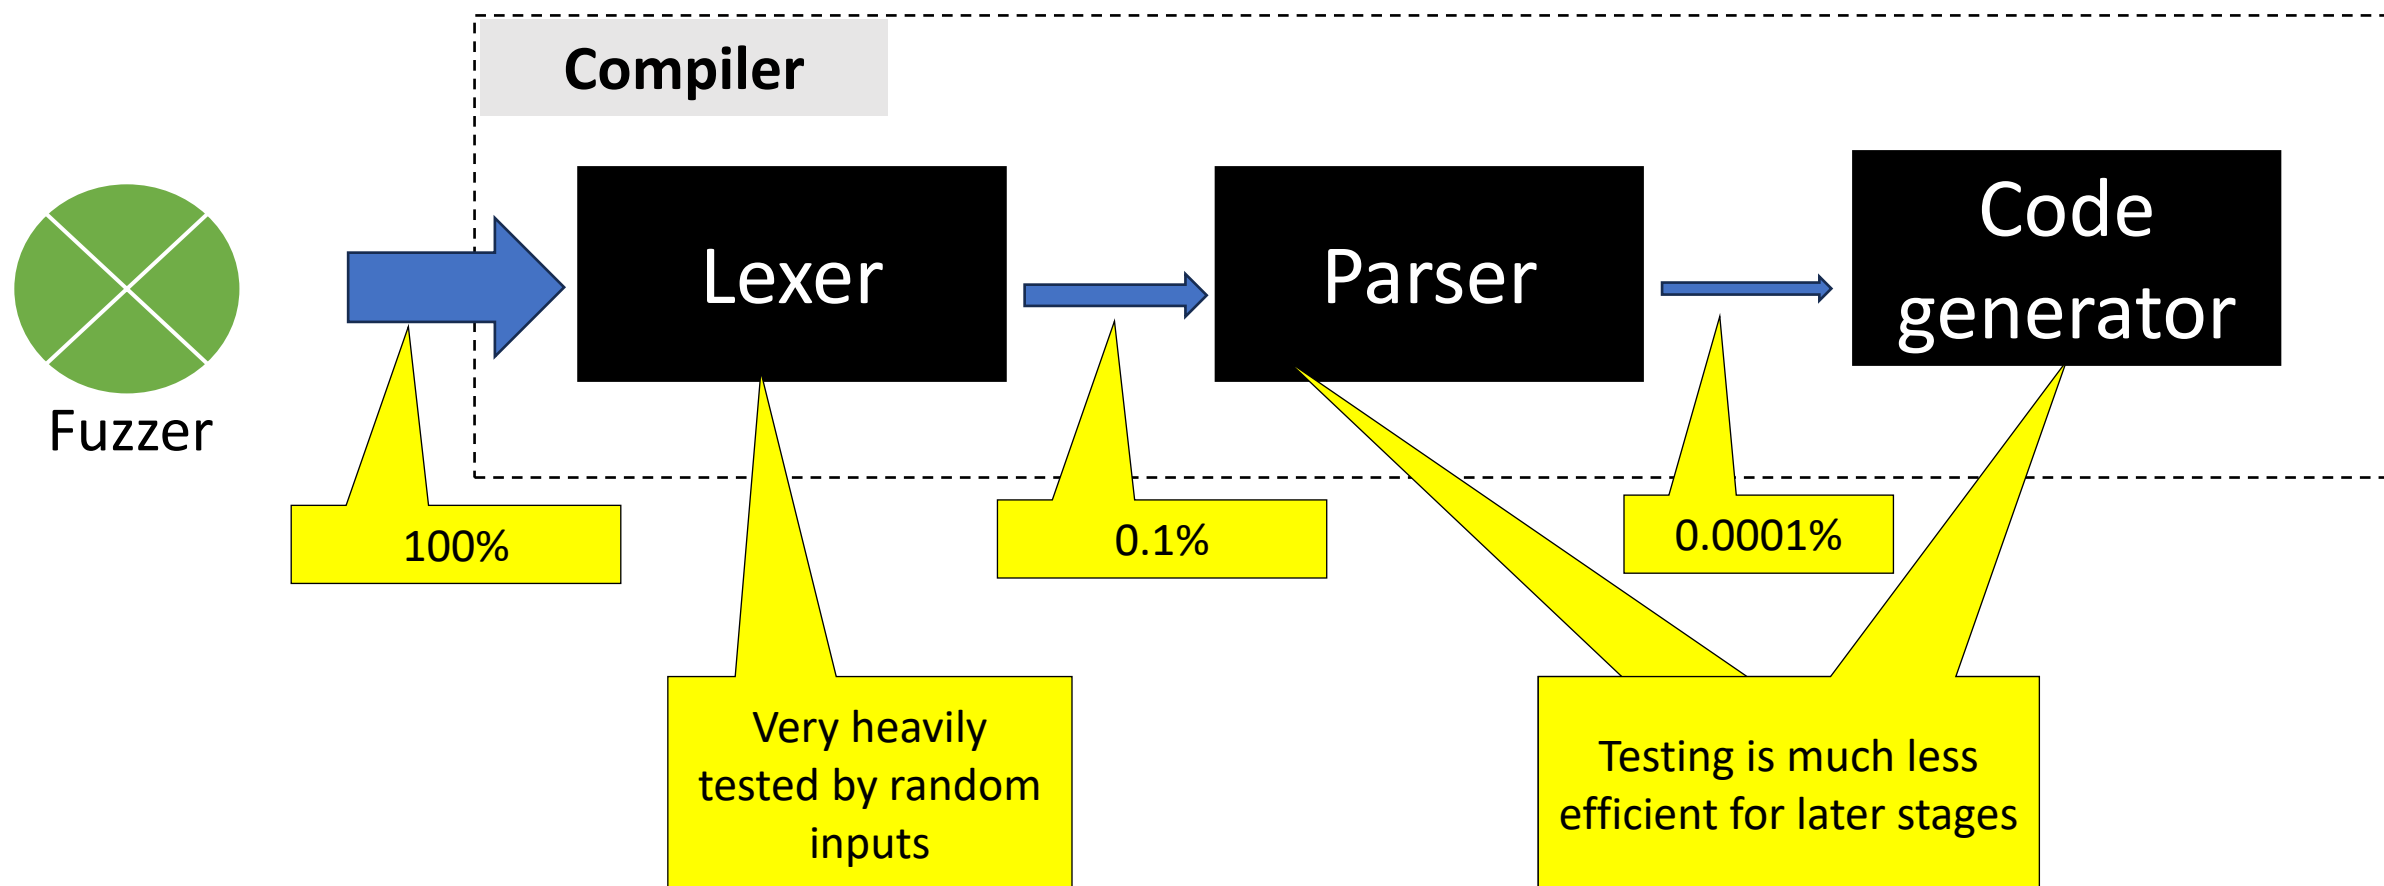
- Cons

- Inefficient test suite
- Might find bugs that are unimportant
- Low code coverage in practice

Uneven Code Coverage: Example 1



Uneven Code Coverage: Example 2



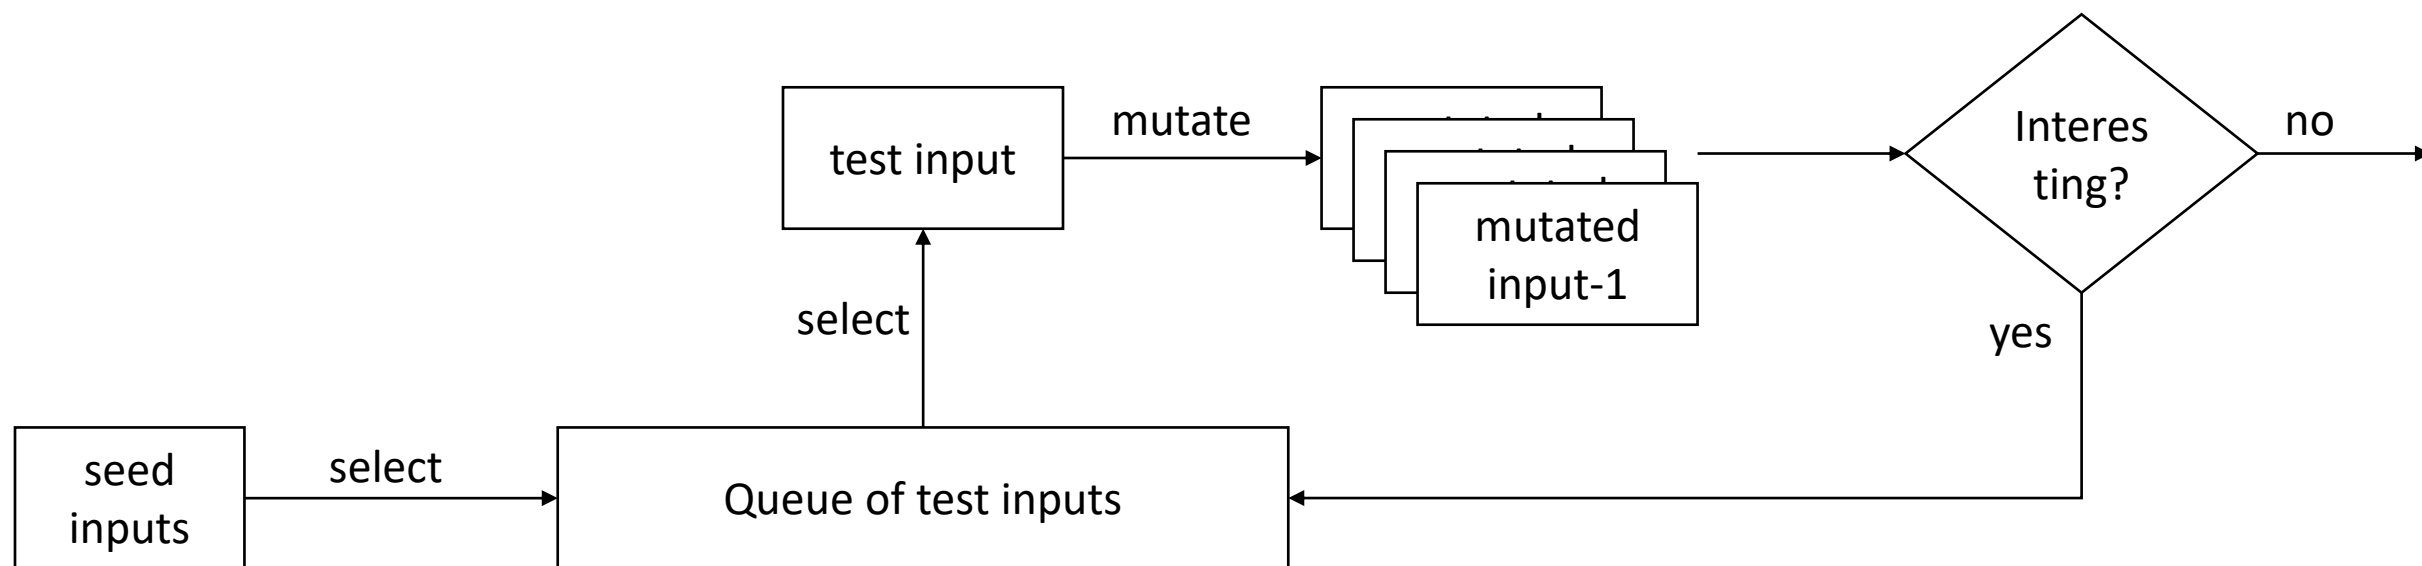
American Fuzzy Lop: AFL

- Gray-box Fuzzing
- Guide input generation toward a **goal**
 - Guidance based on **lightweight program analysis**
- Three main steps
 - **Randomly** generate inputs
 - Get feedback from test executions
 - E.g., what code is **covered**?
 - Mutate inputs that have covered new code

American Fuzzy Lop: AFL



- Simple yet effective fuzzing tool
 - Targets C/C++ programs
 - Inputs are e.g., files read by the program
- Widely used in industry
 - Maintained by Google
 - Find security-related bugs
 - Has found many bugs in OpenSSL, PHP, Firefox, etc.

American Fuzzy Lop: AFL Overview



Control Flow Graphs – Basic Blocks

```
A
if (x>3){
    B
    C
}else{
    D
}
E
F
```



AFL – Measuring Coverage

- AFL uses Branch coverage
 - Branches **between basic blocks**
- Rationale: reaching a code location may not be enough to trigger a bug, but state also matters
- **Tradeoff** between
 - **Effort** spent on measuring coverage
 - **Guidance** it provides to the fuzzer

AFL – Example

Sequence of basic blocks that are
executed

Branches covered (i.e. edges in CFG)

AFL – Example

Sequence of basic blocks that are executed

A -> B -> C -> D -> E

Branches covered (i.e. edges in CFG)

AB, BC, CD, DE

AFL – Example

Sequence of basic blocks that are executed

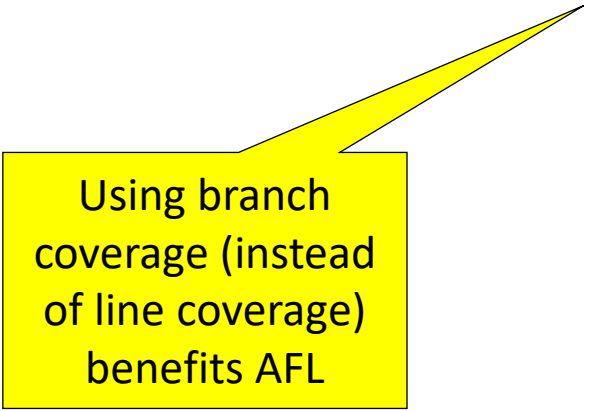
A -> B -> C -> D -> E

A -> B -> D -> C -> E

Branches covered (i.e. edges in CFG)

AB, BC, CD, DE

AB, BD, DC, CE



Using branch coverage (instead of line coverage) benefits AFL

AFL – Efficient Implementation

- AFL instruments all branching points:

```
cur_location = /*RANDOMLY GENERATED AT COMPILE TIME*/;
```

```
shared_mem[cur_location ^ prev_location]++;
```

```
prev_location = cur_location >> 1;
```

To distinguish between
A->B and B->A

Works well with separately
compiled components

Combine current and prev block into
a fixed size hash and increment the
number of times it is seen

AFL – Detecting New Behaviors

- Inputs that **trigger a new edge** in the CFG: considered as **new behavior**
- Alternative: Consider new paths
 - More expensive to track
 - Path explosion problem

AFL – Example

Sequence of basic blocks that are executed

Exec1: A -> B -> C -> D -> E

Branches covered (i.e. edges in CFG)

AB, BC, CD, DE (new => keep input)

AFL – Example

Sequence of basic blocks that are executed

Exec1: A -> B -> C -> D -> E

Exec2: A -> B -> C -> A -> E

Branches covered (i.e. edges in CFG)

AB, BC, CD, DE (new => keep input)

AB, BC, CA, AE (new => keep input)

AFL – Example

Sequence of basic blocks that are executed

Branches covered (i.e. edges in CFG)

Exec1: A -> B -> C -> D -> E

AB, BC, CD, DE (new => keep input)

Exec2: A -> B -> C -> A -> E

AB, BC, CA, AE (new => keep input)

Exec3: A -> B -> C -> A -> B -> C -> D -> E

AB, BC, CA, AB, BC, CD, DE (not new => discard input)

AFL – Example

Sequence of basic blocks that are executed

Exec1: A -> B -> C -> D -> E

Exec2: A -> B -> C -> A -> E

Exec3: A -> B -> C -> A -> B -> C -> D -> E

Indicates a loop. What if new behavior occurs after k iterations of the loop?

Branches covered (i.e. edges in CFG)

AB, BC, CD, DE (new => keep input)

AB, BC, CA, AE (new => keep input)

AB, BC, CA, AB, BC, CD, DE (not new => discard input)

AFL Refinement – Edge Hit Counts

- Refinement of the previous definition of “new behaviors”
- For each edge, count how often it is taken:
 - Approximate counts based on **buckets of increasing size**
 - 1, 2, 3, 4-7, 8-15, 16-31, etc.
 - Rationale: focus on relevant differences in hit counts

AFL – Evolving the Input Queue

- Maintain queue of inputs
 - Initially: **Seed inputs** provided by the user
 - Once used, **keep** input **if it covers a new edge**
 - **Add new inputs by mutating existing input**
- In practice: Queue sizes of 1k to 10k

AFL – Mutation Operators

- Goal: Create new inputs from existing inputs
- Random transformations of bytes in an existing input
 - **Bit flips** with varying lengths and stepovers
 - **Addition and Subtraction** of small integers
 - **Insertion** of known interesting integers
 - E.g., 0, 1, INT_MAX, INT_MIN
 - **Splicing** of different inputs

AFL – More Tricks for Fast Fuzzing

- Time and memory limits
 - Discard input when execution is too expensive
- Pruning the queue
 - Periodically select subset of inputs that still cover every edge seen so far
- Prioritize how many mutants to generate from an input in the queue
 - E.g., focus on unusual paths or try to reach specific locations

AFL – Real World Impact

- Open-source tool maintained mostly by Google
 - Initially created by a single developers
 - Various improvements proposed in academia and industry
- Fuzzers regularly check various security-critical components
 - Many thousands of compute hours
 - Hundreds of detected vulnerabilities

LibFuzzer

- Motivation: enable fuzzing **libraries (software components)** instead of whole application
- User provides fuzzing entry points called fuzz targets
- Intuition: if program has X lines of code and Y fuzz targets, then fuzzer has to cover only X/Y lines of code, on average per target.

LibFuzzer

- Motivation: enable fuzzing **libraries (software components)** instead of whole application
- User provides fuzzing entry points called **fuzz targets**
- Intuition: if program has **X** lines of code and **Y** fuzz targets, then fuzzer has to cover only **X/Y** lines of code, on average per target.

OSS-Fuzz

- Continuous fuzzing infrastructure hosted on Google cloud platform



- OSS-Fuzz has discovered over 17,400 bugs from 2016 – 2019 in many large projects (openssl, llvm, postgresql, git, firefox, ...)

ClusterFuzz

- Google's scalable fuzzing infrastructure used to fuzz Chrome browser
- As of Jan 2019, it has found ~16,000 bugs in Chrome and ~11,000 bugs in over 160 projects integrated with OSS-Fuzz
- Features:
 - Highly scalable (runs over 1000s of machines)
 - Automatically tracks bugs found by underlying Fuzzers
 - Test case minimization
 - Statistics to analyze fuzzers performance
 - Easy to use web interface for reports

Domain-Specific Fuzzing


- Random testing is a paradigm as opposed to a technique
- Two Case Studies:
 - Mobile Apps: Google's Monkey Tool for Android Apps
 - Concurrent programs: Microsoft's Cuzz Tool

Testing Mobile Apps

Android
Framework
Interface

```
class MainActivity extends Activity implements OnClickListener {  
    void onCreate(Bundle bundle) {  
        Button buttons = new Button[] { play, stop, ... };  
        for (Button b : buttons) b.setOnClickListener(this);  
    }  
    void onClick(View target) {  
        switch (target) {  
            case play:  
                startService(new Intent(ACTION_PLAY));  
                break;  
            case stop:  
                startService(new Intent(ACTION_STOP));  
                break;  
            ...  
        }  
    }  
}
```

Function called by
Android
framework
whenever user
taps on any
button




Music player App

Testing Mobile Apps

```
class MainActivity extends Activity implements OnClickListener {  
    void onCreate(Bundle bundle) {  
        Button buttons = new Button[] { play, stop, ... };  
        for (Button b : buttons) b.setOnClickListener(this);  
    }  
    void onClick(View target) {  
        switch (target) {  
            case play:  
                startService(new Intent(ACTION_PLAY));  
                break;  
            case stop:  
                startService(new Intent(ACTION_STOP));  
                break;  
        }  
    }  
}
```

TOUCH(136,351)


TOUCH(136,493)



- TOUCH(x,y) where x and y are randomly generated: x in [0,480] and y in [0,800]

Testing Mobile Apps

```
class MainActivity extends Activity implements OnClickListener {  
    void onCreate(Bundle bundle) {  
        Button buttons = new Button[] { play, stop, ... };  
        for (Button b : buttons) b.setOnClickListener(this);  
    }  
    void onClick(View target) {  
        switch (target) {  
            case play:  
                startService(new Intent(ACTION_PLAY));  
                break;  
            case stop:  
                startService(new Intent(ACTION_STOP));  
                break;  
        }  
    }  
}
```



The diagram illustrates a mobile application interface for a 'Random Music Player'. The screen displays several circular buttons for playback control. Two specific touch events are highlighted with red callout boxes: 'TOUCH(136,351)' pointing to the 'play' button and 'TOUCH(136,493)' pointing to the 'stop' button. These coordinates correspond to the x and y positions of the touch events on the screen.

- TOUCH(x,y) where x and y are randomly generated: x in [0,480] and y in [0,800]
- Monkey tool can generate many other kinds of input events (E.g., key press on keyboard, input from devices' trackball, incoming phone call, change GPS location, ...)

Testing Mobile Apps – Sequence of Events




TOUCH(x1, y1)

Click on Eject button to select the audio file or see the default one



TOUCH(x2, y2)

Dialog box asking user to select/view audio file and play it



TOUCH(x3, y3)

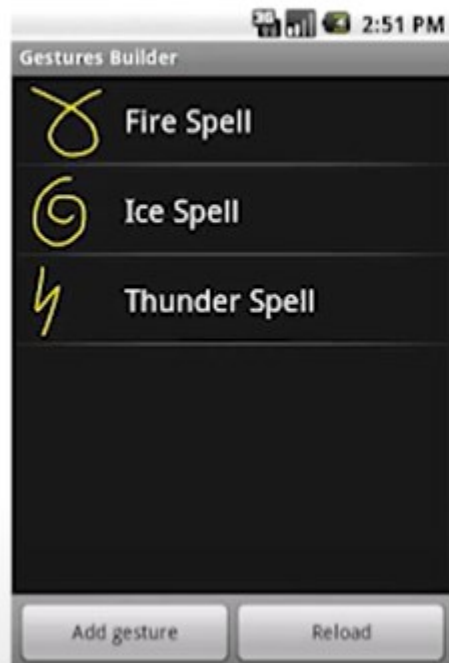
Stop playing the audio file

Testing Mobile Apps – Generating Gestures




Testing Mobile Apps – Generating Gestures

DOWN(x1,y1) MOVE(x2,y2) UP(x2,y2)



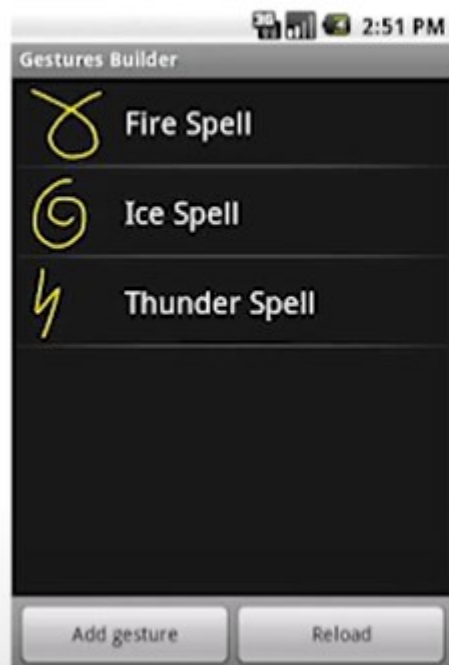
(x1,y1)




(x2,y2)

Testing Mobile Apps – Generating Gestures


DOWN(x1,y1) MOVE(x2,y2) UP(x2,y2)



(x1,y1)



(x2,y2)



Grammar-Based Fuzzing

- Recall: Context Free Grammar
 - Terminals
 - Non-terminals
 - Production rules
 - Start symbol
- Advantages of specifying program inputs using context-free grammars
 - Systematic and Efficient Test Generation
 - Expressive enough to handle complex input formats (e.g., XML, JSON)
 - Provides a basis to fuzz wide range of software components:
 - configurations, APIs, GUIs, protocols, simulations, etc.

Grammar of Monkey Events

test_case := *event* *

event := *action* (*x* , *y*) | ...

action := **DOWN** | **MOVE** | **UP**

x := **0** | **1** | ... | *x_limit*


y := **0** | **1** | ... | *y_limit*

Quiz(1/2): Monkey Events

Use Monkey Grammer to write the event sequences for the specifications

```
test_case := event *  
event := action ( x , y ) | ...  
action := DOWN | MOVE | UP  
x := 0 | 1 | ... | x_limit  
y := 0 | 1 | ... | y_limit
```

Give the specification of a TOUCH event at pixel (80,215).



TOUCH events are a pair of DOWN and UP events at a single place on the screen.

Quiz(2/2): Monkey Events

Use Monkey Grammer to write the event sequences for the specifications

```
test_case := event *
event := action ( x , y ) | ...
action := DOWN | MOVE | UP
x := 0 | 1 | ... | x_limit
y := 0 | 1 | ... | y_limit
```

Give the specification of a TOUCH event at pixel (80,215).

DOWN(80,215) UP(80,215)

TOUCH events are a pair of DOWN and UP events at a single place on the screen.

Give the specification of a MOTION event from pixel (80,215) to pixel (80,100) to pixel (370,100).

MOTION events consist of a DOWN event somewhere on the screen, a sequence of MOVE events, and an UP event.

Quiz(2/2): Monkey Events

Use Monkey Grammer to write the event sequences for the specifications

```
test_case  := event *  
event      := action ( x , y ) | ...  
action     := DOWN | MOVE | UP  
x          := 0 | 1 | ... | x_limit  
y          := 0 | 1 | ... | y_limit
```

Give the specification of a TOUCH event at pixel (80,215).

DOWN(80,215) UP(80,215)

TOUCH events are a pair of DOWN and UP events at a single place on the screen.

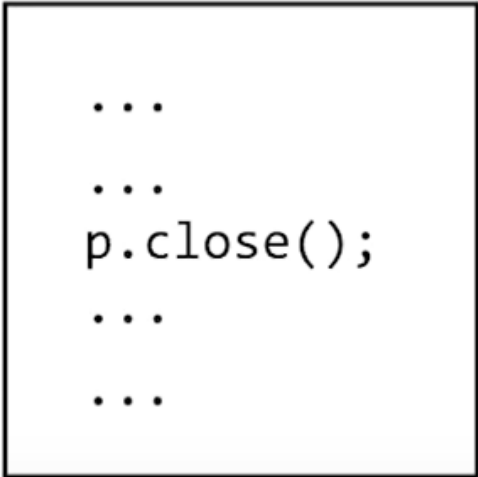
Give the specification of a MOTION event from pixel (80,215) to pixel (80,100) to pixel (370,100).

DOWN(80,215) MOVE(80,100)
MOVE(370,100) UP(370,100)

MOTION events consist of a DOWN event somewhere on the screen, a sequence of MOVE events, and an UP event.

Testing Concurrent Programs

Sequential Program:




```
...  
...  
p.close();  
...  
...
```

Testing Concurrent Programs


Sequential Program:

Input:

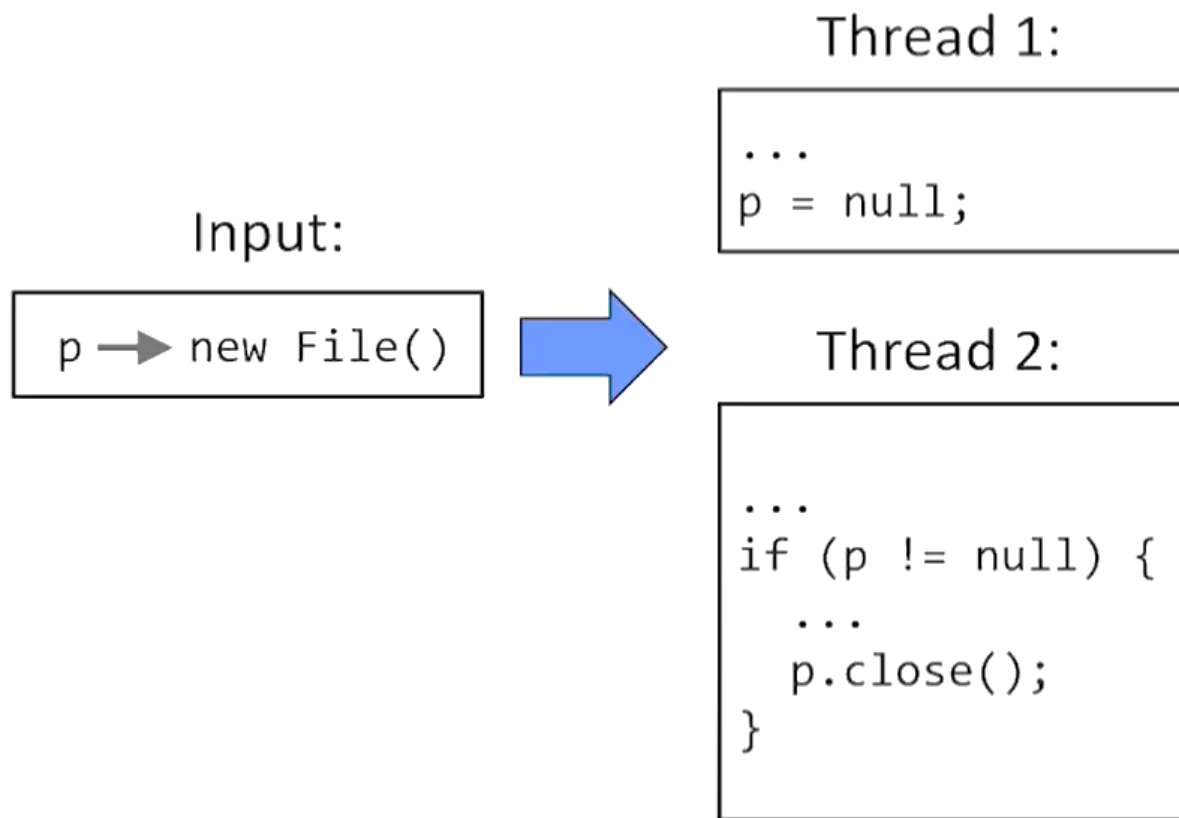
$p \rightarrow \text{null}$



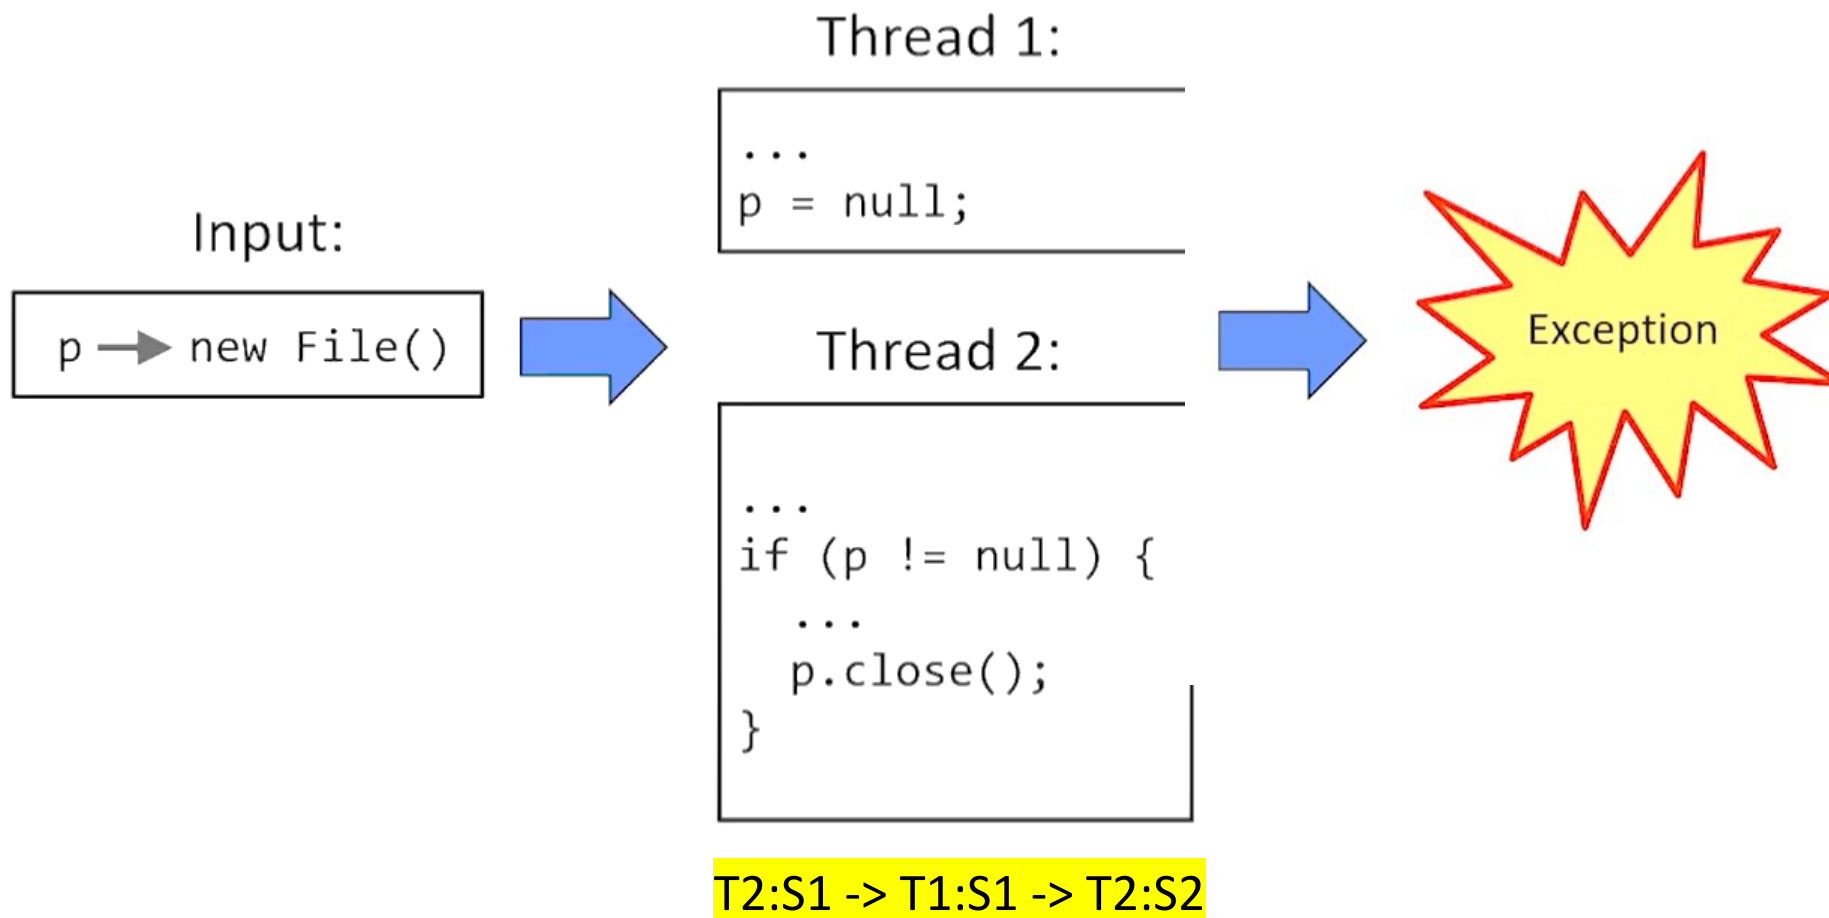
```
...
...
p.close();
...
...
```



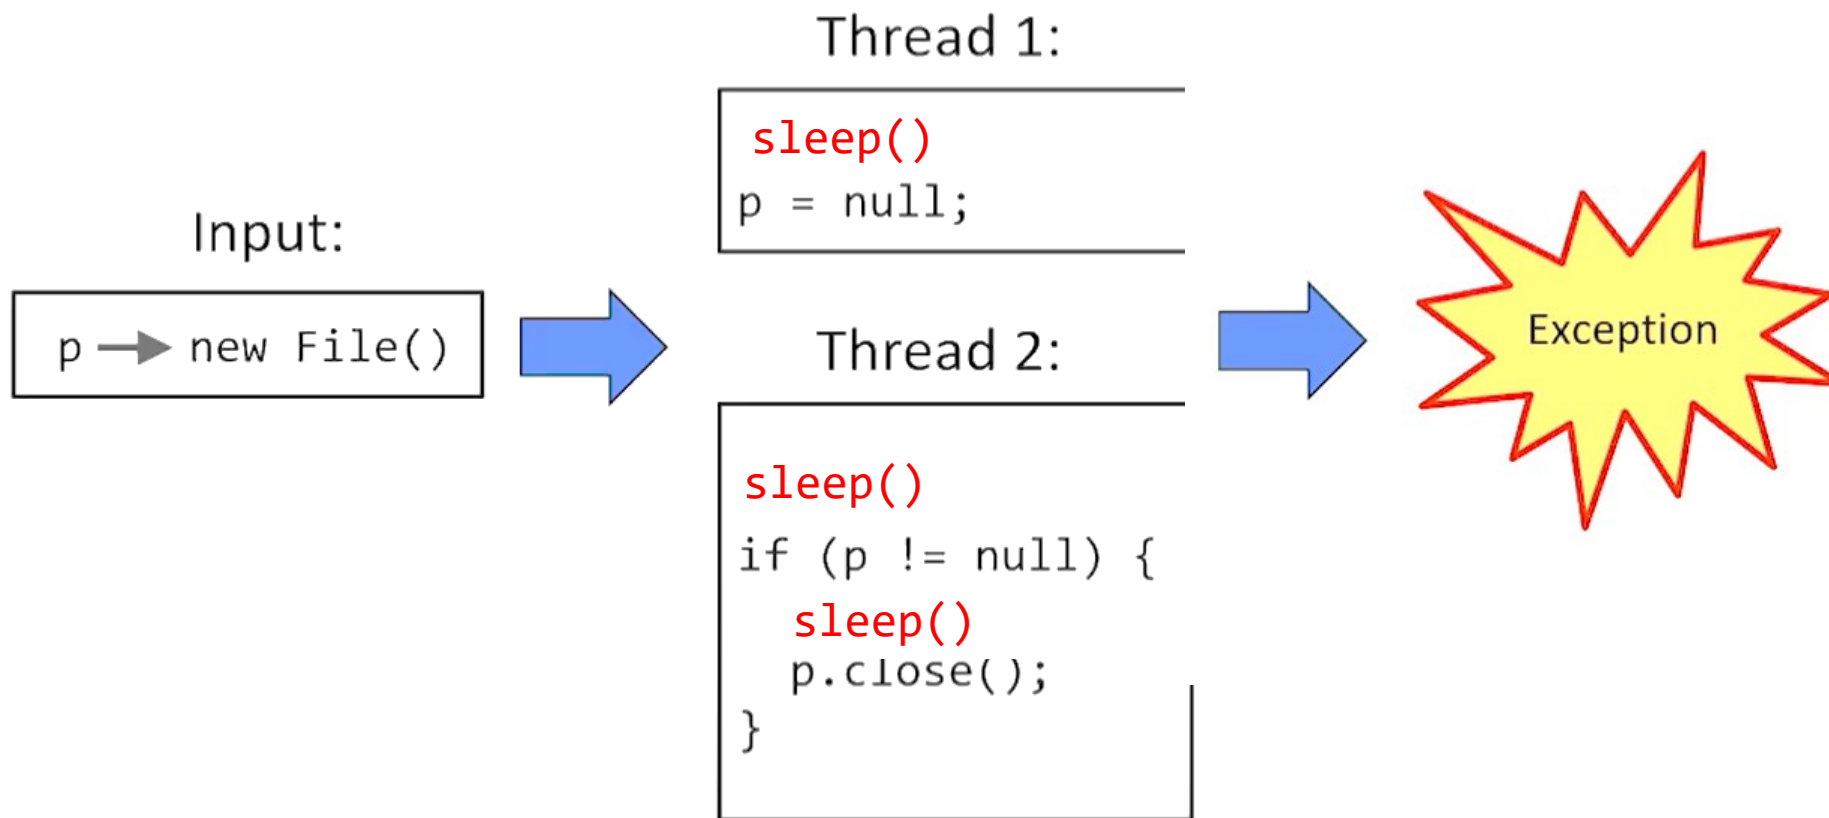
Testing Concurrent Programs



Testing Concurrent Programs



Testing Concurrent Programs



Cuzz: Fuzzing Thread Schedules


- Introduces **sleep()** calls
 - Automatically (instead of manually)
 - Systematically before each statement (as opposed to those chosen by tester)
 - Less tedious and less error-prone
- Gives worst-case probabilistic guarantees on bug finding

Depth of Concurrency Bug

- **Bug depth** = number of **scheduling constraints** that must be satisfied to trigger the bug
- **Scheduling constraint** = requirement on the ordering between two statements in **different threads**.

Depth of Concurrency Bug: Example 1

- **Bug depth** = number of **scheduling constraints** that must be satisfied to trigger the bug
- **Scheduling constraint** = requirement on the ordering between two statements in **different threads**.



Depth of Concurrency Bug: Example 1

- **Bug depth** = number of **scheduling constraints** that must be satisfied to trigger the bug
- **Scheduling constraint** = requirement on the ordering between two statements in **different threads**.

Thread 1:

```
...  
T t = new T();
```


Thread 2:

```
...  
if(t.state == 1)  
...  
...
```

Depth of this concurrency bug is 1


Depth of Concurrency Bug: Example 2

- **Bug depth** = number of **scheduling constraints** that must be satisfied to trigger the bug
- **Scheduling constraint** = requirement on the ordering between two statements in **different threads**.



Depth of Concurrency Bug: Example 2

- **Bug depth** = number of **scheduling constraints** that must be satisfied to trigger the bug
- **Scheduling constraint** = requirement on the ordering between two statements in **different threads**.



Depth of Concurrency Bug

- **Bug depth** = number of **scheduling constraints** that must be satisfied to trigger the bug
- **Scheduling constraint** = requirement on the ordering between two statements in **different threads**.
- Observation exploited by Cuzz: most concurrency bugs typically have small depth
 - “small test case” hypothesis: if there is a bug, there will be some small input that will trigger the bug

Quiz: Concurrency Bug

Specify the depth of the concurrency bug in the following example:

Then specify all ordering constraints needed to trigger the bug. Use notation (x,y) to mean statement x comes before statement y, and separate multiple constraints by a space.

Thread 1

```
1: lock(a);
2: lock(b);
3: g = g + 1;
4: unlock(b);
5: unlock(a);
```

Thread 2

```
6: lock(b);
7: lock(a);
8: g = 0;
9: unlock(a);
10: unlock(b);
```

Quiz: Concurrency Bug

Specify the depth of the concurrency bug in the following example:

2

Then specify all ordering constraints needed to trigger the bug. Use notation (x,y) to mean statement x comes before statement y, and separate multiple constraints by a space.

(1,7) (6,2)

Thread 1

```
1: lock(a);
2: lock(b);
3: g = g + 1;
4: unlock(b);
5: unlock(a);
```

Thread 2

```
6: lock(b);
7: lock(a);
8: g = 0;
9: unlock(a);
10: unlock(b);
```

Deadlock preventing both threads to proceed and makes program hang


Cuzz: Probabilistic Guarantee

- Given a program with:
 - n threads (\sim tens)
 - k steps (\sim millions)
 - Bug depth d (1 or 2)
- Cuzz will find the bug with probability $\geq \frac{1}{nk^{(d-1)}}$ in each run



Worst case guarantee

Cuzz: Probabilistic Guarantee



- Probability (choose correct thread ordering) $\geq 1/2$ (generally $1/n$ for n threads)
- Probability (switch thread ordering at correct statement) $\geq 1/k$, where k = total statements in program
- Probability (triggering bug) $\geq 1/nk$
- For a bug of depth d , thread priorities are changed $d-1$ times, i.e., Cuzz needs to pick $d-1$ statements in program to switch thread execution ordering.
 - Probability (picking right set of $d-1$ statement) $\geq 1/k^{(d-1)}$
- Probability (triggering bug) $\geq 1/nk^{(d-1)}$

What Have We Learnt

- Random Testing is effective is effective for testing mobile apps, security, and concurrency
- It is not a technique but a paradigm that you can apply to solve many different kinds of problems (e.g., Domain-Specific fuzzers)
- Complements and does not replace formal testing
- Must generate inputs from a reasonable distribution to be effective
- May be less effective for systems with multiple layers (e.g. compilers)

Next Class

Delta Debugging

Reminder

- **Paper Presentation** assignment will be released today.
- It will be done in groups of **up to three students**.
- **Project Idea presentation** in next class.
- Each group will give 10 min presentation. The submission for this assignment is due **before** class.
 1. **The Problem.** Tell us what you are going to build. If you are doing a research-focused project, tell us the research question(s) you will answer. If you are building a system, describe the prototype and the basic functionality and where your work will focus.
 2. **The Design.** Tell us how you will build what you are building. If you are doing a research-focused project, tell us the design of your experiment(s). If you are building a system, show up some early design.
 3. **The Evaluation.** Tell us how you will know you succeeded. If you are doing a research-focused project, tell us what data you will use, how you will know that your results make sense, what statistical tests you'll apply, etc. If you are building a system, tell us your testing plan and how you will execute it.
 4. **The Plan.** Tell us your planned timeline and each group member's responsibilities.